

Non-parametric Statistics: Notes 5

Correlation

Gregory Matthews ¹

¹Department of Mathematics and Statistics
Loyola University Chicago

Fall 2014

Outline

- ▶ When you think of correlation you are probably actually thinking of **Pearson correlation**.
- ▶ Pearson correlation measure the linear association between two variables.

$$\rho = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{\sqrt{\sum_{i=1}^n (X_i - \bar{X})^2 \sum_{i=1}^n (Y_i - \bar{Y})^2}} = \frac{\text{cov}(X, Y)}{\sigma_X \sigma_Y}$$




Figure : Pearson Correlation








Figure : Pearson Correlation: $\rho = 0.816$ in each graph

```
plot(x,y)
```



```
#Calucalte the Pearson correlation coefficient using funct  
(rho<-cor(x,y))
```

```
## [1] 0.2382
```

```
#Correlation test
```

```
#What are we testing here?
```

```
cor.test(x,y)
```

```
##
```

```
## Pearson's product-moment correlation
```

```
##
```

```
## data: x and y
```

```
## t = 1.041, df = 18, p-value = 0.3118
```

```
## alternative hypothesis: true correlation is not equal to
```

```
## 95 percent confidence interval:
```

```
## -0.2283 0.6158
```

```
## sample estimates:
```

```
## cor
```


```
## 0.2382
```

```
#permutation test!  
nsim<-1000  
rhoVec<-rep(NA,nsim)  
for (i in 1:nsim){  
  yPerm<-sample(y,length(y),replace=FALSE)  
  rhoVec[i]<-cor(x,yPerm)  
}
```



```
hist(rhoVec,xlim=c(-1,1))  
abline(v=rho,col="red",lwd=5)
```

Histogram of rhoVec




```
#two sided p-value
sum(rhoVec>=abs(rho))/nsim+sum(rhoVec<=-abs(rho))/nsim

## [1] 0.301

cor.test(x,y)

##
## Pearson's product-moment correlation
##
## data:  x and y
## t = 1.041, df = 18, p-value = 0.3118
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
##  -0.2283  0.6158
## sample estimates:
##      cor
## 0.2382
```

- ▶ What if there is a relationship between two variables, but the relationship is not linear?
- ▶ For instance, what is the relationship is strictly monotonic?
- ▶ Let's take a look at that situation.



#If we look at the Pearson correlation we get:

```
cor(x,y)
```

```
## [1] 0.9059
```

#But in some sense these two variables are perfectly correlated

#What if we computed correlation based on the ranks?

```
xR<-rank(x);yR<-rank(y)
```

```
cor(xR,yR)
```

```
## [1] 1
```

#This is Spearman correlations

```
cor(x,y,method="spearman")
```

```
## [1] 1
```

- ▶ Spearman correlation is Pearson correlation using ranks of the data.
- ▶ Spearman correlation measure the strength of the monotonic relationship between two variables.

$$\rho = \frac{\sum_{i=1}^n (R_i^X - \bar{R}^X)(R_i^Y - \bar{R}^Y)}{\sqrt{\sum_{i=1}^n (R_i^X - \bar{R}^X)^2 \sum_{i=1}^n (R_i^Y - \bar{R}^Y)^2}} = \frac{\text{cov}(R^X, R^Y)}{\sigma_{R^X} \sigma_{R^Y}}$$

#Test the null hypothesis that the Spearman correlation is

#In other words there is no association between X and Y

```
cor.test(x,y,method="spearman")
```

```
##
```

```
## Spearman's rank correlation rho
```

```
##
```

```
## data: x and y
```

```
## S = 0, p-value = 0.01667
```

```
## alternative hypothesis: true rho is not equal to 0
```

```
## sample estimates:
```

```
## rho
```

```
## 1
```

#This p-value of this test is equal to

```
2*(1/120)
```

```
## [1] 0.01667
```

```

library(gtools)
perms<-permutations(5,5)
corNull<-rep(NA,dim(perms)[1])
for (i in 1:dim(perms)[1]){
  corNull[i]<-cor(x,y[perms[i,]],method="spearman")
}
sort(corNull)

```

```

##      [1] -1.0 -0.9 -0.9 -0.9 -0.9 -0.8 -0.8 -0.8 -0.7 -0.7
##     [15] -0.6 -0.6 -0.6 -0.6 -0.6 -0.6 -0.6 -0.5 -0.5 -0.5
##     [29] -0.4 -0.4 -0.4 -0.3 -0.3 -0.3 -0.3 -0.3 -0.3 -0.3
##     [43] -0.2 -0.2 -0.2 -0.2 -0.2 -0.1 -0.1 -0.1 -0.1 -0.1
##     [57] -0.1  0.0  0.0  0.0  0.0  0.0  0.0  0.1  0.1  0.1
##     [71]  0.1  0.1  0.1  0.2  0.2  0.2  0.2  0.2  0.2  0.3
##     [85]  0.3  0.3  0.3  0.3  0.3  0.4  0.4  0.4  0.4  0.5
##     [99]  0.5  0.6  0.6  0.6  0.6  0.6  0.6  0.6  0.7  0.7
##    [113]  0.8  0.8  0.8  0.9  0.9  0.9  0.9  0.9  1.0

```


Two quick notes

- ▶ When there are ties, simply take the average of the ranks and proceed as usual.
- ▶ Careful using Spearman correlation (or any correlation). If observations are correlated, Spearman may result in incorrect results. **Only use when observations are uncorrelated.**

Kendall's τ

- ▶ First we need to talk about **concordance** and **discordance**.
- ▶ Two pairs of data (X_i, Y_i) and (X_j, Y_j) are concordance if $X_i < X_j$ implies that $Y_i < Y_j$.
- ▶ Two pairs of data (X_i, Y_i) and (X_j, Y_j) are discordance if $X_i < X_j$ implies that $Y_i > Y_j$.
- ▶ Concordance means that the changes in the X's and Y's are in the same direction.

Kendall's τ

- ▶ If pairs of observations are more likely to be concordant, then we say there is a positive relationship.
- ▶ If pairs of observations are more likely to be discordant, then we say there is a negative relationship.
- ▶ Based on the definition of concordance, we can also express concordance as $(X_i - X_j)(Y_i - Y_j) > 0$. (WHY?)
- ▶ And we can write $\tau = 2P[(X_i - X_j)(Y_i - Y_j) > 0] - 1$
- ▶ $P[(X_i - X_j)(Y_i - Y_j) > 0]$ is the probability of concordance. Then we rescale this number to range from -1 to 1 like the other measures of correlation.

Kendall's τ

- How do we estimate Kendall's τ (assuming no ties.)

Define:

$$U_{ij} = \begin{cases} 1 & : (X_i - X_j)(Y_i - Y_j) > 0 \\ 0 & : (X_i - X_j)(Y_i - Y_j) < 0 \end{cases}$$

Then let:

$$V_i = \sum_{j=i+1}^n U_{ij}$$

V_i is the number of pairs that are concordant with the i -th pair for $j \geq i + 1$ Then:

$$r_\tau = 2 \frac{\sum_{i=1}^{n-1} V_i}{\binom{n}{2}} - 1$$

```
#let's compute Kendall's Tau by hand (ONCE!)  
(dat<-data.frame(x=c(1,2,3,4),y=c(2,4,3,1)))
```

```
##      x y  
## 1 1 2  
## 2 2 4  
## 3 3 3  
## 4 4 1
```

```
#Concordant pairs (1,2),(1,3)  
#Disconcordant pairs (1,4),(2,3),(2,4),(3,4)  
#2*(numConcordPairs/TotalNumberOfPairs)-1  
2*(2/6)-1
```

```
## [1] -0.3333
```

```
#Or just use the R function  
cor(dat$x,dat$y,method="kendall")
```

```
## [1] -0.3333
```

```
#Null hypothesis that there is no association  
#Versus alternative that there is some association  
cor.test(dat$x, dat$y, method="kendall")  
  
##  
## Kendall's rank correlation tau  
##  
## data: dat$x and dat$y  
## T = 2, p-value = 0.75  
## alternative hypothesis: true tau is not equal to 0  
## sample estimates:  
## tau  
## -0.3333
```

Contingency Tables

- ▶ When we compare two categorical variables to each other, we often use two way contingency tables.
- ▶ We say there are R levels of the first variable and C levels of the second variable, then we can create an R by C table.
- ▶ The elements of the contingency table are counts of the individuals who fall into both categories corresponding to a cell in the table.

```
#A contingency table
```

```
dat<-data.frame(x=c(1,1,1,2,2,2,3,3,3,4,4,4),y=c(1,3,2,4,2,  
table(dat)
```

```
##      y  
## x    1 2 3 4  
##    1 1 1 1 0  
##    2 0 1 0 2  
##    3 1 0 1 1  
##    4 1 1 1 0
```


Contingency Tables

- ▶ What hypothesis are we interested in testing here? It depends!
- ▶ Consider two cases:
 1. All n individuals are selected at random and cross classified according to row and column characteristic.
 2. A fixed number $n_{i\cdot}$ for $i = 1, 2, \dots, r$ are sampled and classified according to column characteristic.

- ▶ First define: $p_{ij} = \frac{E[n_{ij}]}{n}$
- ▶ Case 1: $H_0 : p_{ij} = p_{i.}p_{.j}$ vs $H_a : p_{ij} \neq p_{i.}p_{.j}$
- ▶ Case 2: $H_0 : p_{j|i} = p_{j|i'}$ vs $H_a : p_{j|i} \neq p_{j|i'}$ where $p_{j|i} = \frac{p_{ij}}{p_{i.}}$
- ▶ The null hypotheses in both Case 1 and Case 2 are equivalent EVEN THOUGH the sampling situations are different.

- ▶ How do we test these hypotheses?
- ▶ Use the following test statistic. (Why is this a good test statistic?)

Test Statistic:

$$\chi^2 = \sum_{i=1}^r \sum_{j=1}^c \frac{(n_{ij} - e_{ij})^2}{e_{ij}}$$

where $e_{ij} = \frac{n_{i.} n_{.j}}{n}$

- ▶ If all $e_{ij} > 5$, the $\chi^2 \sim \chi^2_{df=(r-1)(c-1)}$
- ▶ However, when expected cell frequencies are small, the approximation may not be appropriate.
- ▶ This is especially true when many cells have no responses.
- ▶ What can we do in that situation? Permutation test!

#Table 5.4.2 in the book

```
dat<-data.frame(who=factor(c(rep("Physician-Prescribed",4)),  
table(dat)
```

##		satisfaction		
##	who	Not	Somewhat	Very
##	Physician-Prescribed	2	2	0
##	Self-Administered	0	1	2

```
(XsqObs<-chisq.test(table(dat))$statistic)
```

```
## Warning:  Chi-squared approximation may be  
incorrect
```

```
## X-squared
```

```
##      4.278
```

```
library(gtools)
perms<-combinations(7,4)
XsqPerms<-rep(NA,35)
for (i in 1:dim(perms)[1]){
  datTemp<-dat
  ind<-c(perms[i,],setdiff(c(1:7),perms[i,]) )
  datTemp$satisfaction<-datTemp$satisfaction[ind]
  XsqPerms[i]<-chisq.test(table(datTemp))$statistic
}
```

```
## Warning: Chi-squared approximation may be
incorrect
```

```
## Warning: Chi-squared approximation may be
incorrect
```

```
## Warning: Chi-squared approximation may be
incorrect
```

```
## Warning: Chi-squared approximation may be
incorrect
```

```
## Warning: Chi-squared approximation may be
```

```
XsqObs
```

```
## X-squared
```

```
##      4.278
```

```
table(round(XsqPerms,2))
```

```
##
```

```
## 0.19 2.24 4.28 4.96      7
```

```
##   12   12    6    4    1
```

```
sum(XsqPerms>=XsqObs)/35
```

```
## [1] 0.3143
```


- ▶ Real quick. The special case of the permutation test for a 2 by 2 contingency table is called the Fisher's exact test.
- ▶ This is implemented in R using the function `fisher.test`
- ▶ Example follows:

Agresti (1990, p. 61f; 2002, p. 91) Fisher's Tea Drinker
A British woman claimed to be able to distinguish whether milk or
tea was added to the cup first. To test, she was given 8 cups of
tea, in four of which milk was added first. The null hypothesis
is that there is no association between the true order of pouring
and the woman's guess, the alternative that there is a positive
association (that the odds ratio is greater than 1).

```
TeaTasting <-  
matrix(c(3, 1, 1, 3),  
       nrow = 2,  
       dimnames = list(Guess = c("Milk", "Tea"),  
                        Truth = c("Milk", "Tea")))
```

TeaTasting

```
##      Truth  
## Guess  Milk Tea  
##  Milk    3   1  
##  Tea     1   3
```

#Why do we want one-sided here?

```
fisher.test(TeaTasting, alternative = "greater")
```

```
##
```

```
## Fisher's Exact Test for Count Data
```

```
##
```

```
## data: TeaTasting
```

```
## p-value = 0.2429
```

```
## alternative hypothesis: true odds ratio is greater than
```

```
## 95 percent confidence interval:
```

```
## 0.3136 Inf
```

```
## sample estimates:
```

```
## odds ratio
```

```
## 6.408
```

=> $p = 0.2429$, association could not be established

McNemar's test

- ▶ McNemar's test is used for paired comparison studies when the responses are dichotomous.
- ▶ The classic example of this is to ask people which candidate they prefer before a political debate, and then ask them again after the debate.
- ▶ If there are two candidates (A and B) there are four possibilities: (A,A) (A,B) (B,A) (B,B).
- ▶ We are interested in testing the null hypothesis that the probability of A before is the same as the probability of A after.
- ▶ $H_0 : P_{AA} + P_{AB} = P_{AA} + P_{BA} = P_{AB} = P_{BA}$

McNemar's test

- ▶ If we let $n = X_{AB} + X_{BA}$ and $P_{AB} = P_{BA}$, then there is an equal chance of switching from A to B or B to A.
- ▶ In this way the distribution of $X_{AB}|n \sim \text{binomial}$ with mean $0.5n$ and variance $0.25n$. WHY?
- ▶ We can then use this to calculate a p-value.

```
## Agresti (1990), p. 350.
## Presidential Approval Ratings.
## Approval of the President's performance in office in two surveys,
## one month apart, for a random sample of 1600 voting-age Americans.
Performance <-
matrix(c(794, 86, 150, 570),
       nrow = 2,
       dimnames = list("1st Survey" = c("Approve", "Disapprove"),
                        "2nd Survey" = c("Approve", "Disapprove")))
Performance
```

```
##           2nd Survey
## 1st Survey Approve Disapprove
## Approve      794      150
## Disapprove   86      570
```

```
(n<-Performance[1,2]+Performance[2,1])
```

```
## [1] 236
```

```
#calculate the p-value
```

```
2*(1-pbinom(Performance[1,2]-1,n,.5))
```

```
## [1] 3.716e-05
```

```
2*(pbinom(Performance[2,1],n,.5))
```

```
## [1] 3.716e-05
```



```
#approximate the p-value
T4<-((Performance[1,2]-0.5*n)/sqrt(0.25*n))^2
1-pchisq(T4,1)

## [1] 3.099e-05

mcnemar.test(Performance,correct=FALSE)

##
## McNemar's Chi-squared test
##
## data: Performance
## McNemar's chi-squared = 17.36, df = 1, p-value = 3.099e-05
```

```
mcnemar.test(Performance,correct=TRUE)
```

```
##
```

```
## McNemar's Chi-squared test with continuity correction
```

```
##
```

```
## data: Performance
```

```
## McNemar's chi-squared = 16.82, df = 1, p-value = 4.115e-05
```

```
## => significant change (in fact, drop) in approval ratings
```